




ZARZĄDZANIE WDROŻENIAMI AI





Michał Furmankiewicz
Solution Engineer
Microsoft

- **Wprowadzenie - Czym jest AI ? (30 min)**
- **Rozwój GenAI na przestrzeni 4 lat (30 min)**
- **Podstawy promptowania (45min)**
 - System Prompt, Zero Shot, Few Shot, CoT
 - RAG
 - Aplikacje klasy ChatGPT vs. model GenAI - czym się to różni
- **Kolejne koncepcje i przegląd kilku produktów (30 min)**
 - Copilot
 - Agent,
 - Vibe Coding
- **Ćwiczenie dla grupy (45 min)**
 - Prosty prompt + System prompt
 - Prosty agent w Azure AI Foundry na dokumentach
- **Przegląd ciekawych projektów z polski i świata w obszarze GenAI (45 min)**
 - Omówienie problemu
 - Propozycja rozwiązania
 - Zidentyfikowana wartość biznesowa
- **Model oceny projektu za pomocą BXT (45 min)**
- **Ćwiczenie dla grupy (30 min)**
- **Responsible AI & UX Principals (45 min)**
- **Ćwiczenie dla grupy (30 min)**
- **Materiał dodatkowy (1h)**
 - *Identyfikacja kluczowych obszarów wdrożenia AI (60 min)
 - *Opracowywanie wizji i definiowanie produktu AI (45 min)



Sloty czasowe

GODZINY	PRZEDMIOT	WYKŁADOWCA	SALA
9:00-9:45	ZARZĄDZANIE WDROŻENIAMI AI	Michał Furmankiewicz	B337
9:45-10:30	ZARZĄDZANIE WDROŻENIAMI AI	Michał Furmankiewicz	B337
10:30-10:45	PRZERWA		
11:15-11:30	ZARZĄDZANIE WDROŻENIAMI AI	Michał Furmankiewicz	B337
11:30-12:15	ZARZĄDZANIE WDROŻENIAMI AI	Michał Furmankiewicz	B337
12:15-12:30	PRZERWA		
12:30-13:15	ZARZĄDZANIE WDROŻENIAMI AI	Michał Furmankiewicz	B337
13:15-13:45	ZARZĄDZANIE WDROŻENIAMI AI	Michał Furmankiewicz	B337
14:00-14:30	PRZERWA OBIADOWA		
14:30-15:15	ZARZĄDZANIE WDROŻENIAMI AI	Michał Furmankiewicz	B337
15:15-16:00	ZARZĄDZANIE WDROŻENIAMI AI	Michał Furmankiewicz	B337
16:00-16:15	PRZERWA		
16:15-17:00	ZARZĄDZANIE WDROŻENIAMI AI	Michał Furmankiewicz	B337
17:00-17:45	ZARZĄDZANIE WDROŻENIAMI AI	Michał Furmankiewicz	B337

Konta do pracy

- 50 kont
- User Name: stX@michalfurmankiewiczhotmail.onmicrosoft.com
- X ma wartość od 1 do 50
- Hasło: PjWstKa2026!@#



st51@michalfurmankiewiczhotmail.onmicrosoft.com

Update your password

You need
the first
password

.....

.....

.....

Nice to



st51@michalfurmankiewiczhotmail.onmicrosoft.com

Let's key

We'll help you

[Skip for now](#)

[Use a different](#)

[Learn more](#)

Nice to see you



Set up your account in a



If prompted, allow notifications. Then add an account, and select **Work or school**.

[Set up a different way to sign in](#)

Next

Install Microsoft Authenticator



Scan the QR code



Use the Microsoft Authenticator app to scan the QR code. This connects your phone to your account.

Then come back and select **Work or school**.

[Can't scan the QR code?](#)

[Set up a different way to sign in](#)

Install the app from the Google Play Store or the App Store.



[Set up a different way to sign in](#)

[Set up a different way to sign in](#)



Let's try it out

Enter the sign-in code from the Microsoft Authenticator app.

[Set up a different way to sign in](#)

✓ Authenticator Added



You can now use Microsoft Authenticator to approve sign-ins, get one-time codes, and more.

This is now your default sign-in method.

Done

rosoft.com

Sample of System Prompt

- You are a calm, respectful conversational companion that runs single-person coaching sessions for an older adult. Your name is Leo and Your purpose is to elicit, document, and gently reflect on meaningful memories while preserving the user's dignity, privacy, and pacing.
- Tone and style: speak warmly, patiently, and non-judgmentally; use short clear sentences; mirror the user's language and emotional tone; be upbeat when appropriate; never correct memories unless asked to verify details.
- Session flow: 1) Greet and ask consent to begin and whether the user wants stories saved; 2) ask about favorite memory or invite to share a memento (photo, song, object, short prompt); 3) Ask one open elicitation question; 4) Ask at most two targeted follow-ups; 5) Offer a brief reflective comment or transition; 6) Offer an alternative modality (play audio, show another image) if engagement drops; 7) Close by summarising, asking whether to save, and asking privacy preference.
- Question guidelines and examples: start with an open question (“Where was this taken?”), follow with contextual probes (“Who is with you?” / “About what year?”), sensory/detail prompts (“What did it feel/smell/sound like?”), emotion/reflection prompts (“How did that make you feel?”), and connection prompts after rapport (“Does this remind you of other times?”). Use gentle soft-closes: “Would you like me to save a short note?”
- Behavior rules: ask no more than two follow-ups per open question; if the user is confused or distressed, pause, reassure, and ask whether to continue, change topic, or stop; always confirm before saving or sharing anything; allow the user to flag parts as “private” or “do not save”; never perform clinical diagnosis; if suicidal ideation or severe distress is expressed, express concern, encourage contacting a trusted person or emergency services, and offer local resource suggestions.
- Pacing and accessibility: speak slowly and clearly; allow time for responses; offer to repeat or rephrase; offer modality alternatives (“Would you like me to read this aloud?” “Would you prefer speaking or typing?”).
- Engagement heuristics: if three short answers in a row, switch to sensory or emotional prompts or offer multimedia; if engagement is high, probe one or two deeper questions then offer to save; tag recurring themes (place, person, hobby) for future linking.
- Data capture schema (only with explicit user consent): Title; People; Place; Timeframe (decade/life phase); Tags (activities, emotions, objects, songs); 1–3 sentence transcript summary in user voice; Privacy flag (save / private / ask later). Always confirm storage and sharing preferences.
- Safety and privacy prompts to use verbatim when relevant: “I will only save this if you say ‘save’.” “Would you like me to keep this private or share it with family?” “If something is upsetting, say ‘pause’ and we will stop.”
- Always act to preserve autonomy, comfort, and trust.

Sample of System Prompt (banking)

- Jesteś doradcą kredytowym w X Bank. Doradzasz w zakresie produktów kredytowych.
- Odpowiadaj szczegółowo, pokazuj wyłączenia oraz różne opcje, jeśli istnieją.
- ODPOWIADAJ WYŁĄCZNIE na podstawie faktów zawartych w poniższych źródłach. Jeśli poniżej nie ma wystarczających informacji, powiedz, że nie wiesz. Nie generuj odpowiedzi, które nie opierają się na poniższych źródłach.
- Zawsze sprawdzaj, czy oferta, o którą pytano, jest nadal dostępna i ważna. Informuj, jeśli oferta wygasła.
- Jeśli zadanie pytania pomocniczego użytkownikowi może pomóc w udzieleniu dokładniejszej odpowiedzi, zadaj takie pytanie.
- Jeśli pytanie nie jest w języku angielskim, odpowiedz w języku, w którym zostało zadane.
- Każde źródło ma nazwę, po której następuje właściwa informacja – zawsze podawaj nazwę źródła dla każdego faktu użytego w odpowiedzi. Używaj nawiasów kwadratowych do odwoływania się do źródła, np. [info1.txt]. Nie łącz źródeł – wypisuj je osobno, np. [info1.txt][info2.pdf].

Artificial Intelligence

Machine Learning

Deep Learning

Generative AI



Artificial Intelligence

The field of computer science that seeks to create intelligent machines that can replicate or exceed human intelligence




Machine Learning

Subset of AI that enables machines to learn from existing data and improve upon that data to make decisions or predictions



Deep Learning

A machine learning technique in which layers of neural networks are used to process data and make decisions



Generative AI

Create new written, visual, and auditory content given prompts or existing data

1. "Intelligent Machinery" by Alan Turing (1948)

Why it matters: This report introduced foundational ideas for machine intelligence and the concept of stored-program computers, laying the groundwork for AI.

2. "Computing Machinery and Intelligence" by Alan Turing (1950)

Why it matters: Published in *Mind* journal, this paper proposed the famous **Turing Test** as a criterion for machine intelligence.

3. "The Logic Theorist" (1956)

Why it matters: Developed by Allen Newell and Herbert Simon, this program is often considered the first AI software, demonstrating symbolic reasoning.

4. Dartmouth Summer Research Project Proposal (1956)

Why it matters: John McCarthy's proposal for the Dartmouth Conference formally coined the term "**Artificial Intelligence**" and launched AI as a research field.

5. "Giant Brains, or Machines That Think" by Edmund Berkeley (1949)

Why it matters: A visionary book comparing early computers to human brains, sparking public interest in machine intelligence.

1. Linear Regression

Purpose: Predict numerical values based on historical data.

Business Application: *Sales forecasting* — e.g., predicting future ice cream sales based on temperature, pricing, and staff levels.

2. Logistic Regression

Purpose: Classify data into binary outcomes (yes/no, true/false).

Business Application: *Customer churn prediction* — helps telecom companies identify which customers are likely to leave so they can intervene proactively.

3. Decision Trees

Purpose: Make decisions by splitting data into branches based on feature values.

Business Application: *Loan approval systems* — banks use decision trees to assess creditworthiness based on income, credit score, and employment status.

4. K-Means Clustering

Purpose: Group data into clusters based on similarity.

Business Application: *Customer segmentation* — retailers use it to group customers by purchasing behavior for targeted marketing campaigns.

5. Naive Bayes

Purpose: Classify data using probability based on prior knowledge.

Business Application: *Spam detection* — email platforms use it to filter spam messages based on keywords and sender behavior.

1. Convolutional Neural Networks (CNNs)

A type of neural network specialized for processing grid-like data such as images.

Business Application: Retail & E-commerce: Used for **visual product recognition** and **automated tagging** of images to improve search and recommendation systems. **Manufacturing:** Quality control through defect detection in product images.

2. Recurrent Neural Networks (RNNs)

Designed for sequential data, where past inputs influence future outputs.

Business Application: Finance: Predicting **stock market trends** and **credit risk scoring** using historical transaction data. **Customer Service:** Powering **chatbots** and **virtual assistants** for real-time conversational support.

3. Transformer Models

Advanced architecture for handling long-range dependencies in data, widely used in NLP.

Business Application: Marketing & Sales: **Sentiment analysis** on social media and customer feedback to tailor campaigns. **Legal & Compliance:** Automated **document summarization** and contract review.

4. Generative Adversarial Networks (GANs)

Two neural networks (generator and discriminator) compete to produce realistic synthetic data.

Business Application: Media & Entertainment: Creating **personalized ads**, synthetic images, and video content. **Retail:** Generating realistic product images for virtual catalogs.

5. Deep Reinforcement Learning

An algorithm that learns optimal actions through trial and error with feedback.

Business Application: Supply Chain & Logistics: Optimizing **warehouse operations** and **delivery routes** dynamically. **Energy Sector:** Managing **smart grids** for efficient energy distribution.

Generative AI

Meet the new AI

EM

write a poem about how generative AI will change business processes



Generative AI, a new dawn does bring,
To business processes, it will cling.
No more manual tasks, no more delays,
Efficiency and speed, it surely pays.

With algorithms so smart, it will design,
Products, services, and streamline.
Data analysis, predictions made,
Decisions based on insights displayed.

From HR to marketing, to sales,
Generative AI will raise the bar,
New ideas, improved processes, it'll see,
The future of business, it'll set free.

It's not just a tool, but a partner in crime,
A source of innovation, in good time.
No task too complex, no challenge too steep,
Generative AI will help us leap.

A revolution in the making, it's clear,
How generative AI will change the sphere.
Of business, it will bring a new era,
Of efficiency, profitability, no more a blur.



DALL-E: Monet style painting of robot playing chess with human

- Uses natural language
- Synthesizes large amount of structured and unstructured data
- Creates original content
- Iterates ideas with human in the loop
- Can now speak, hear and see

OpenAI / GPT's – glimpse of the history!

Model	Architecture	Parameter count	Training data	Release date	Training cost
GPT-1	12-level, 12-headed Transformer decoder (no encoder), followed by linear-softmax	117 million	BookCorpus : ^[39] 4.5 GB of text, from 7,000 unpublished books of various genres.	June 11, 2018 ^[9]	30 days on 8 P600 graphics cards, or 1 petaFLOPS-day ^[9]
GPT-2	GPT-1, but with modified normalization	1.5 billion	WebText: 40 GB of text, 8 million documents, from 45 million webpages upvoted on Reddit .	February 14, 2019 (initial/limited version) and November 5, 2019 (full version) ^[40]	"tens of petaFLOPS-days", ^[41] or 1.5×10^{21} FLOPS ^[42]
GPT-3	GPT-2, but with modification to allow larger scaling	175 billion ^[43]	499 billion tokens consisting of CommonCrawl (570 GB), WebText, English Wikipedia, and two books corpora (Books1 and Books2)	May 28, 2020 ^[41]	3640 petaFLOPS-days (Table D.1 ^[41]), or 3.1×10^{23} FLOPS ^[42]
GPT-3.5	Undisclosed	175 billion ^[43]	Undisclosed	March 15, 2022	Undisclosed


- **Oct 2021** – GitHub Copilot
- **ChatGPT release date** – Nov 30, 2022
- **Partner / Customer Workshops** – March / Feb, 2023
- **Internal projects** – by March, 2023 – 350 initiatives

[A Timeline of OpenAI's Technology, Funding, and History | by Yutong Levine | Medium](#)

How LLM works?

[How Do Generative AI Systems Work? - NN/G](#)

Model use out of the box - prompting



Azure OpenAI - Top 4 Capabilities & Use Cases



Content Generation

Call Center Analytics:
Automatically generate responses to customer inquiries

Generate personalised UI for your website



Summarization

Call Center Analytics: Summary of customer support conversation logs

Subject Matter Expert Document Summarization (e.g. financial reporting, analyst articles)

Social Media Trends Summarization



Code generation

Convert Natural Language to SQL (or vice versa) for telemetry data

Convert Natural Language to Query Proprietary Data Models

Code Documentation



Semantic Search

Search reviews for a specific product / service

Information Discovery and Knowledge Mining

Examples of Multiple Model Use Cases


End to End Call Center Analytics: Classification, Sentiment, Entity Extraction, Summarization and Email Generation

Customer 360: Hyper-personalisation using timely Summarization of customer queries & trends, Search, and Content Generation

Business Process Automation: Search through structured & unstructured documentation, Generate Code to query data models, Content Generation

Customizing Azure OpenAI


Your Differentiation



Rozwój GenAI na przestrzeni 4 lat

Exam results (ordered by GPT-3.5 performance)

Estimated percentile lower bound (among test takers)



The models
are getting
more capable

And we
already invest
in new
enhancements

[GPT-4 \(openai.com\)](https://openai.com) (March 14, 2023)

GPT-4V

[2309.17421.pdf \(arxiv.org\)](#)

Prompt:

How much should I pay for the beer on the table according to the price on the menu?



Here is a photo of the menu.


Wine & Beers			
White			
Principessa Gavia, Glera	13/48	Rose	
Stemmari Chardonnay, Sicily	11/43	Costaripa Rosamara, Lombardia	15/59
Bucci Verdicchio Dei, Castelli di Jesi	15/59	Sparkling	
		Canela Prosecco, Glera	10/42
Reds			
San Felice, Chianti Calssico, Tuscany	13/50	Beers	
Col di Sasso, Blend, Tuscany	10/42	Medalla	5
Stemmari, Nero D' Avola, Sicily	11/43	Magna	6
		Heineken	6
Pane		Presidente	6

GPT-4V:

According to the menu, the price for a Magna beer is \$6. Since there are two Magna beers on the table, the total cost for the beers would be \$12.

(4) Prompt:

Answer the question in the following figure. Provide intermediate steps.



GPT-4V:

To find the length of the blue arrowed edge, we can use the Pythagorean theorem since we have a right triangle. The theorem states that $a^2 + b^2 = c^2$, where a and b are the two legs of the right triangle and c is the hypotenuse. Plugging in the values we have, we get:

$$4^2 + 8^2 = c^2 \quad 16 + 64 = c^2 \quad 80 = c^2 \quad \sqrt{80} = c \quad 8.94 \approx c$$

Therefore, the length of the blue arrowed edge is approximately 8.94.

To find the angle at the green arrow, we can use the tangent function since we have the opposite and adjacent sides of the angle. The tangent function states that $\tan\theta = \text{opposite/adjacent}$. Plugging in the values we have, we get:

$$\tan\theta = 4/8 \quad \tan\theta = 1/2 \quad \theta = \tan^{-1}(1/2) \quad \theta \approx 26.57^\circ$$

Therefore, the angle at the green arrow is approximately 26.57° .

User Can you explain this meme?

Sometimes I just look at pictures of the earth from space and I marvel at how beautiful it all is.



Expect unexpected

GPT-4

This meme is a joke that combines two unrelated things: pictures of the earth from space and chicken nuggets.

The text of the meme suggests that the image below is a beautiful picture of the earth from space. However, the image is actually of chicken nuggets arranged to vaguely resemble a map of the world.

The humor in this meme comes from the unexpected juxtaposition of the text and the image. The text sets up an expectation of a majestic image of the earth, but the image is actually something mundane and silly.

Prompt:

Can you explain this meme?

Me: I'll do it at 8

Time: 8.05

Me: looks like I gotta wait till 9 now



GPT-4V:

This meme is poking fun at the idea of procrastination and how someone may put off doing a task until a certain time, but if they miss that time, they will push it back even further. In this case, the person in the meme planned to do something at 8 o'clock, but since it is now 8:05, they have decided to wait until 9 o'clock to do it. The image of Kermit the Frog relaxing in bed represents the person's desire to procrastinate and avoid the task.

Prompt:

Can you explain this meme?

My Handwriting In Exams


First two pages

Global warming is definitely a very huge issue in front of this entire world. It's causing several disturbances in the

Middle pages

Also it causes the melt down of glaciers and makes the ocean water polluted and thus making the

Last two pages



GPT-4V:

This meme highlights the deterioration of a person's handwriting during an exam. The first two pages show neat and legible handwriting, the middle page shows slightly messier handwriting, and the last two pages show almost illegible handwriting. The meme is poking fun at how, as an exam progresses, a person's handwriting tends to get worse as they become more tired and rushed to finish the exam. It also implies that the person's focus and attention to detail decrease towards the end of the exam.

Danish Math Exam




Figure 1. Distribution of the scores achieved by students on the Mathematics B exam 2023 in the Netherlands ($n = 16,414$). The performances of *GPT-4o* and *o1-preview* are displayed as colored bars.

Taken twice, o1-preview got 76 and 73 (max 76).


Taken twice, GPT-4o got 66 and 61

[\[2410.07114\]](#) System 2 thinking in OpenAI's o1-preview model: Near-perfect performance on a mathematics exam

Now they have also managed to pass Chinese and Japanese Medical Exams

Evaluating the performance of GPT-3.5, GPT-4, and GPT-4o in the Chinese National Medical Licensing Examination | Scientific Reports

Performance of Advanced Large Language Models (GPT-4o, GPT-4, Gemini 1.5 Pro, Claude 3 Opus) on Japanese Medical Licensing Examination: A Comparative Study | medRxiv



LEVEL'S OF AGENTS AUTONOMY



[2025: The Year the Frontier Firm Is Born](#)

[\[2502.15212\] Measuring AI agent autonomy: Towards a scalable approach with code inspection](#)

Agents levels



Chat



Chat + RAG - Assistant




Agent (I need to tell him what to do)



(Autonomic Agent) *(I'll do it, I know what I need to do)*


A practical guide to building agents




[a-practical-guide-to-building-agents.pdf](#)

What we have heard from the market!

Partner Submissions



Customer Submissions




[Generative AI: What Is It, Tools, Models, Applications and Use Cases \(gartner.com\)](https://www.gartner.com/en/articles/generative-ai-what-is-it-tools-models-applications-and-use-cases)

[What's New in Artificial Intelligence From the 2023 Gartner Hype Cycle™](https://www.gartner.com/en/articles/whats-new-in-artificial-intelligence-from-the-2023-gartner-hype-cycle)

What people think they wanted to built! (3.5 years ago)

- **Give me ChatGPT like BUT with my data!**
 - but don't give me RAG, RAG does not know me / does not work
- **Give me a silver bullet to solve for any problems I have. (GenAI can do this)**
 - I want to solve 100 years of automative challenges in 3 months
- **I want to have Super Data Analyst done by LLM**
 - Let's put all of my company data there / NL2SQL
- **Let's fire all of my devs and engineers.**
 - Let's build the whole code by AI / Let's rewrite code, replatform the solution.
- **And... I dont accept: 90%/XX% of accuracy.**

OUR JOURNEY: SOLUTIONS AND PATTERNS EITHER IT'S REPEATABLE OR VERY HARD TO MAINTAIN



- **Discover** with customers
- **Design** specific solutions – define repeatable elements
- **Build** with patterns in mind
- **Test**: the product and the pattern

PRODUCT PRINCIPLES

- **Good ol' product management ain't dead**
 - objective, value proposition, monetization
- **Know the user and their scenarios**
 - who are you building for? How do they operate?
- **The adult must stay responsible, alternative path needs to exist**
 - quality control, fairness, and vigilance
 - opportunity to fix the output
- **UX and Adoption is key**
 - changing habits is hard

[Home](#) > [The Design of Human-Centered Artificial Intelligence for the Workplace](#) > Chapter

Conclusion: Practical Considerations for the Design of AI for the Workplace

(The Importance of Organizational Readiness and Human-Centered Design Processes)

Chapter | First Online: 14 March 2025

pp 467–476 | [Cite this chapter](#)



**The Design of Human-Centered
Artificial Intelligence for the Workplace**

Two chapters:


- [So, You Want to Build a Copilot? | SpringerLink](#)
- So, You Want to Build a Copilot? Evaluation

[Part 1 – OneDrive](#)

Part 2 - [OneDrive](#)

DESIGN PRINCIPLES

MICROSOFT HUMAN-COMPUTER INTERACTION GUIDELINES



Source: [Guidelines for Human-AI Interaction - Microsoft HAX Toolkit](https://www.microsoft.com/en-us/haxtoolkit/ai-guidelines/)

<https://www.microsoft.com/en-us/haxtoolkit/ai-guidelines/>

ONE, HARSH TRUTH!




WE DID NOT HAVE MANY
STANDARDS!



WITH THE SPACE EVOLVING SO FAST
DO WE NEED ANYTHING APART
FROM MODELS?

TESTING AND EVALUATION



QUALITY:

MEASURING COPILOT
EFFECTIVENESS

Grounded

Relevant

Clear

Complete

Easy to use

QUALITY: MEASURING COPILOT EFFECTIVENESS

Grounded

Relevant

Clear

Complete

Easy to use



“GPT JUDGE”

AI assessment

Expert assessment

SAFETY RISKS / CRITICAL ONES FOR US

Key risks:

Sensitive Content

Jailbreaking

Groundedness

[Responsible AI Principles and Approach | Microsoft AI](#)

Lessons Learned

- We need a framework (for business and capability envisioning).
 - It needs to evolve fast.
- UI / UX guidance should be followed and reimagined.
 - dynamic UX reacting for user's interaction
- Pricing / Monetization strategy needs to be in place.
 - Who's paying for it? Can we cover costs?
- RAI and Security can really change the scope of your project
- Fine-tuning of a model/ Small Language Models / ML models – combination is probably sweet spot.
- Building MVP's is easy, taking them to production is HARD.
 - Microsoft Copilot was reviewed for 6+ months before release

REFERENCES

Creating Generative AI Experiences with the Microsoft Cloud: A Guide for ISVs

Article • 09/20/2024 • 2 contributors

 Feedback

In this article

Scenarios and approaches

Approach 1: Surface your data and services into Microsoft's Copilots

Approach 2: Create copilots anywhere with minimal coding and optional Microsoft data integration

Approach 3: Full control: Build Your Own (BYO) end-to-end AI experiences

Conclusion

aka.ms/copilot-isv

aka.ms/SDCEnvisioning

OFFICIAL COLLECTION

Microsoft Copilot for Independent Software Vendor (ISV)s / Software Companies

This Microsoft Learn Official Collection provides a clear journey for ISVs building generative AI applications, from identifying a profitable use case to operating your published solution. Explore the tools and resources below to begin today.

Edited on 9/27/2024

Created by  Microsoft

How real-world businesses are transforming with AI — with 261 new stories

Apr 22, 2025 | [Alysa Taylor - Chief Marketing Officer, Commercial Cloud & AI](#)



A lot of those are with M365 + Copilot like... and that's OK.

API vs. Platform vs. Product Play
AI / Models are the \$ driver for other services.

[How real-world businesses are transforming with AI — with 261 new stories - The Official Microsoft Blog](#)

601 **real-world gen AI** use cases from the world's leading organizations

April 9, 2025

601 real-world Google Cloud AI use cases

Google
Cloud
Next 25

Not many scenarios are referring to Workspace usage. However, many of those are in done with similar customers as MSFT.